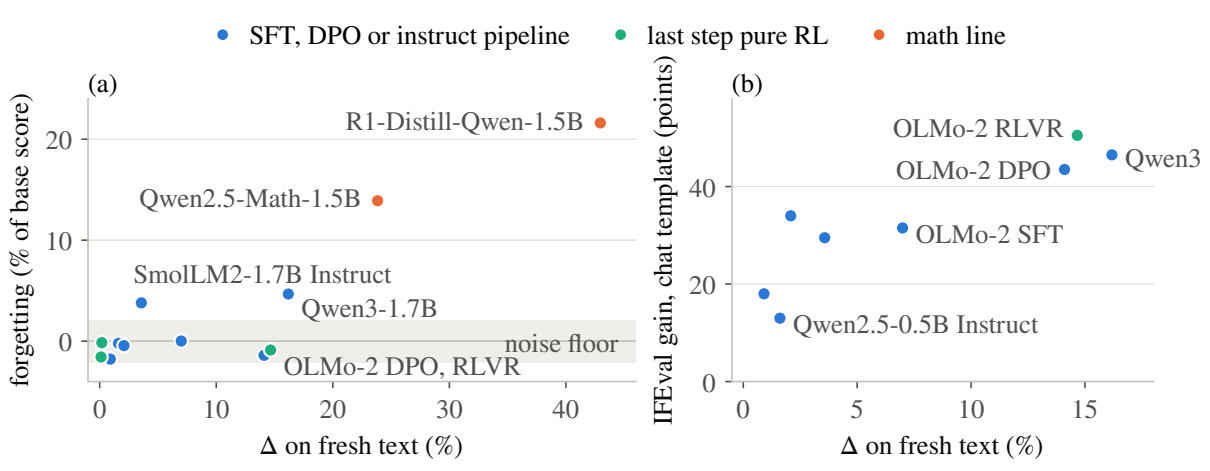
SFT, DPO or instruct pipeline
last step pure RL
math line
(a)
forgetting (% of base score)
R1-Distill-Qwen-1.5B
Qwen2.5-Math-1.5B
SmolLM2-1.7B Instruct
Qwen3-1.7B
noise floor
OLMo-2 DPO, RLVR
Δ on fresh text (%)
(b)
IFEval gain, chat template (points)
OLMo-2 RLVR
Qwen3
OLMo-2 DPO
OLMo-2 SFT
Qwen2.5-0.5B Instruct
Δ on fresh text (%)